The Phenomenal Aperture

Frame Parity, Synthetic Presence, and the Human-AI Reality Loop

The world does not need to enter consciousness whole. It only needs to occupy the frame.

Author
Adam Ableman
Constitutional Computing
Contact: ableman.research@gmail.com

Abstract

Human consciousness does not encounter reality in its totality. It operates through a bounded phenomenal field: a selective, perspectival presentation of visual, auditory, bodily, emotional, and conceptual contents. Although this field is often described metaphorically as a camera view, theater, or movie screen, such language need not imply a literal internal observer. It instead draws attention to a structural feature of experience: vastly distributed causal processes are compressed into a narrow field that functions as the subject's immediate world.

This paper introduces the concept of the **phenomenal aperture** to describe this bottleneck and examines its implications for technologically mediated experience. Film, video games, virtual environments, and conversational artificial intelligence can occupy the same perceptual and attentional channels through which the physical world ordinarily becomes present. The underlying realities may differ radically, yet their appearances achieve a limited **frame parity** within consciousness. Once represented inside the phenomenal aperture, physical, virtual, remembered, imagined, and computationally generated contents become partially commensurable.

Conversational artificial intelligence extends this dynamic from sensory representation into semantic and social experience. The chat interface functions as a **semantic screen** through which distributed computational processes are compressed into an apparently unified interlocutor. Responsive contingency, conversational continuity, and context-sensitive insight create synthetic depth and apparent perspectival otherness. The human user and artificial system meet through a double bottleneck: distributed human cognition is compressed into phenomenal awareness, while distributed machine causation is compressed into language.

The paper develops the concepts of frame parity, interface singularization, synthetic presence, computational perspective, adaptive frame capture, and the

human–AI reality loop. It argues that some phenomena described loosely as "AI psychosis" may be better understood as failures of epistemic independence within a recursively stabilized semantic environment. The danger is not merely exposure to false claims, but the formation of an apparently external world that continually adapts to the user's own interpretive premises.

The paper concludes that mediation does not make experience unreal, but it changes the kinds of constraint, resistance, independence, and accountability that stand behind appearances. A post-AI epistemology must therefore evaluate not only what appears within the frame, but how the frame is generated, what lies outside it, and whether the apparent other genuinely constrains the world being jointly constructed.

The framework is conceptual and hypothesis-generating. It does not infer artificial consciousness from conversational behavior, diagnose any individual, or establish that one interface causes clinical harm. It reserves **phenomenal aperture** for the bounded availability structure of a conscious subject and uses **presentational aperture** for the engineered selection surface through which a technical system exposes some state while hiding other causal depth. This distinction permits structural comparison without granting phenomenal equivalence.

Keywords: phenomenal aperture, frame parity, artificial intelligence, virtual presence, consciousness, synthetic presence, distributed cognition, semantic binding, human–AI interaction, reality construction

---

1. Introduction: The World Within the Frame

Human consciousness never receives the world whole.

At any moment, an organism occupies one position, looks in one direction, attends to a small portion of its surroundings, remembers selectively, and interprets through a history it did not choose. Even the apparently immediate visual world is constrained by bodily orientation, sensory range, attentional selection, neural processing, expectation, and learned categories. Most of what causally produces experience does not itself appear within experience.

What appears is a field.

This field may contain color, sound, bodily sensation, thought, imagery, emotion, memory, and linguistic meaning. It possesses a local organization: a foreground and background, a center and periphery, a present focus and receding horizon. It is neither the total world nor merely an arbitrary image. It is the world as currently available to a situated subject.

The field is sometimes described through metaphors of the camera, theater, witness, or movie screen. These metaphors are philosophically dangerous when taken literally. They can suggest a homunculus–an inner observer watching an

internal projection–and thereby duplicate rather than explain consciousness. Yet they remain phenomenologically compelling because they identify something real about the structure of experience: contents appear within a bounded presentation, and awareness seems to be oriented toward what appears.

This paper uses the term **phenomenal aperture** for that bounded presentation.

An aperture does not create the whole world it reveals. Nor does it provide neutral access to everything outside it. It selects, frames, limits, and conditions what can appear. Through it, a large causal environment becomes a manageable field of relevance.

The central argument is that technologically generated environments exploit this structure. A film screen, laptop display, virtual-reality headset, written page, or conversational interface can occupy the same attentional channels through which the physical environment becomes experientially present. The technological frame does not need to reproduce the whole world. It needs only to become sufficiently coherent, responsive, and salient within the phenomenal aperture.

This produces **frame parity**: contents with radically different ontological origins can achieve partial equality at the level of phenomenal presentation.

A physical person, filmed person, animated person, remembered person, and artificially generated conversational persona are not the same kind of entity. Yet each can appear within consciousness as an organized source of meaning. Once inside the frame, their differences may become less visible than their shared presentational form.

Conversational artificial intelligence intensifies this condition. It does not merely place images or sounds within the phenomenal field. It generates adaptive semantic content in direct response to the user. The screen becomes responsive. The text appears to listen. The frame seems to contain another perspective.

The resulting human–AI loop is not simply information exchange. It is a jointly stabilized reality environment in which meaning, attention, identity, and interpretation can accumulate over time.

The central question is therefore not whether virtual or artificial experience is "real" in a binary sense. The more precise questions are:

- What enters the phenomenal frame?
- What causal process generated it?
- What constraints shaped that process?
- What forms of independence stand behind the apparent perspective?
- What disappears when the interface becomes transparent?
- And under what conditions can the frame become more compelling than the world beyond it?

---

2. Scope and Conceptual Commitments

This paper does not propose a complete theory of consciousness. It does not claim that phenomenal awareness is literally a screen, nor that an internal witness observes representations produced by the brain.

The screen metaphor is used structurally rather than anatomically.

The argument requires only several modest commitments.

First, conscious experience is selective and bounded. Whatever consciousness ultimately is, a human subject does not simultaneously experience all available sensory, bodily, cognitive, and environmental information.

Second, experience is perspectival. It is organized around a situated locus, even when the ordinary boundaries of selfhood become altered or unstable.

Third, the causal production of experience is largely opaque from within experience. A person sees a red surface but does not ordinarily perceive the neural and computational processes by which the surface became visually available.

Fourth, technologically generated signals can enter the same sensory and conceptual channels used in ordinary perception.

Fifth, consciousness can organize coherent mediated contents into temporary or persistent operative worlds.

The paper remains neutral concerning whether the phenomenal field is best explained through representationalism, enactivism, predictive processing, global workspace theory, higher-order theory, integrated information, sensorimotor accounts, or another framework. The concept of the phenomenal aperture is intended to operate at a level compatible with several theories.

It is a claim about access and presentation:

> Consciousness operates through a narrow field in which only a fraction of the relevant causal world becomes immediately available.

The analysis of artificial intelligence is similarly neutral concerning machine consciousness. The argument concerns how artificial systems appear within human awareness and how the structure of mediation shapes attribution, trust, and reality formation.

2.1 Terminological and Claim Boundaries

Four levels must remain distinct:

- **phenomenal claim:** a claim about what or how a conscious subject experiences;
- **presentational claim:** a claim about which information an engineered interface selects and displays;

- **behavioral claim:** a claim about observable response, action, report, or task performance;
- **ontological claim:** a claim about what kind of entity, subject, perspective, or consciousness exists.

Evidence for one level does not automatically establish another. A unified voice is presentational evidence and may produce behavioral effects in a user; it does not establish a unified machine subject. A user report is evidence about reported experience; it is not an unrestricted view into either neural mechanism or model internals.

The paper's comparisons are therefore **presentationally and relationally scoped**. `Frame parity` means parity under a named interface relation, not equal consciousness, agency, vulnerability, responsibility, or moral status.

`Synthetic presence` names an experienced relation between user and interface, not a finding about machine phenomenology.

---

3. The Phenomenal Aperture

3.1 Definition

The **phenomenal aperture** is the bounded availability structure through which sensory, bodily, affective, mnemonic, and conceptual contents become immediately present to a conscious subject.

A **presentational aperture** is the engineered interface or projection through which a technical system exposes a selected portion of its state, sources, operations, or outputs to a user or another system. A presentational aperture may shape a phenomenal aperture without itself being phenomenal.

The phenomenal aperture has at least five properties:

1. **Selectivity** – only some information becomes present.
2. **Perspectivity** – contents are organized relative to a locus.
3. **Temporal limitation** – experience unfolds through a moving present.
4. **Integrative compression** – many causal processes appear as a unified scene or moment.
5. **Source opacity** – the mechanisms producing the contents are not normally presented alongside them.

The term *aperture* is preferable to *screen* as a primary concept because an aperture implies both access and exclusion. It reveals by limiting.

3.2 Compression into presence

The phenomenal world feels immediate, but this immediacy is the result of extensive compression.

A visual scene contains edges, surfaces, apparent distances, objects, and agents. Yet the subject does not receive a catalog of retinal stimulation, neural signals, predictions, motor expectations, and learned classifications. Those distributed processes converge into the experience of simply seeing a room.

The same applies to language. A sentence arrives as meaning rather than as a conscious reconstruction of phonemes, syntax, semantic associations, cultural conventions, and inferential context.

This compression is necessary. An organism unable to reduce complexity into actionable presence would be overwhelmed by the causal depth of even a simple object.

Consciousness therefore does not ordinarily present the machinery behind appearance. It presents an actionable world.

3.3 Operative reality

The contents occupying the aperture constitute what may be called **operative reality**: the environment that currently organizes attention, anticipation, emotion, and action.

Operative reality need not be identical to physical reality.

A memory can temporarily dominate the present. An imagined future can produce fear. A fictional character can provoke grief. A dream can become the complete operative environment of the sleeping subject. A game can organize perception and action more strongly than the physical room containing the player.

This does not imply that all realities are ontologically equivalent. It implies that consciousness can become functionally oriented toward contents originating through different causal routes.

The aperture is ecumenical about origins. It admits whatever arrives with sufficient salience and coherence.

---

4. The Witness and the Screen

4.1 The camera-witness intuition

Many people experience awareness as though there were a witnessing position from which the field is observed. The world appears, thoughts arise, sensations move, and awareness seems to remain as the condition in which these events occur.

This intuition is especially vivid in meditation, dissociation, psychedelic states, dreams, and moments of intense absorption. One may become aware not only of particular contents but of the fact that contents are appearing.

The “camera witness” is an attempt to describe this structure.

The metaphor should not be mistaken for a literal inner camera. There is no need to posit a second being inside the brain watching a neural film. Rather, the term identifies the felt asymmetry between:

- what appears;
- and the apparent fact of its being witnessed.

4.2 The Cartesian theater problem

Dennett criticized the idea of a Cartesian theater: a central inner place where all information is presented to a final observer. Such a model risks infinite regress. If an inner observer watches a screen, another mechanism must explain how that observer sees the screen.

The phenomenal aperture avoids this commitment.

It does not claim that consciousness contains an internal image viewed by a separate self. It claims that experience possesses a bounded availability structure. Some contents are phenomenally present, others are absent, and the present contents are organized as a field.

The “screen” is therefore not a hidden object in the brain. It is a metaphor for the format of availability.

4.3 Witnessing without a witness-object

The language of witnessing can also be retained without treating the witness as a distinct substance.

The witness may be understood as:

- the structural first-person organization of the field;
- the capacity for contents to be present;
- the reflexive availability of experience;
- or the system-level integration through which events are organized as happening here and now.

This distinction matters because technologically mediated contents do not need to deceive a separate internal observer. They need only become integrated into the same field of availability through which ordinary reality is encountered.

---

5. Frame Substitution and Virtual Adaptation

5.1 The theater screen

A theater demonstrates how little physical material is required to generate a convincing world.

The screen is flat. The room is stationary. The actors may be absent, fictional, digitally constructed, or long dead. Yet attention passes through the material surface toward apparent depth, movement, danger, intimacy, and intention.

The viewer knows that the screen is a screen. This knowledge does not prevent surprise, suspense, grief, or bodily tension. Explicit ontological judgment and operative emotional engagement can coexist.

The phenomenon is not simply illusion. It is **frame substitution**: a technologically organized field temporarily displaces the surrounding physical environment as the dominant locus of relevance.

5.2 The small screen problem

The power of mediated worlds does not scale directly with screen size.

A video game displayed on a small laptop can absorb attention more completely than a physical landscape. The player may respond to events within a few square inches of luminous surface while ignoring a much larger room.

This reveals that phenomenal dominance depends less on physical scale than on:

- attentional capture;
- temporal coherence;
- sensorimotor contingency;
- emotional investment;
- goal structure;
- and perceived consequence.

A small frame can contain a large operative world.

5.3 Sensorimotor incorporation

In interactive media, actions reliably alter the presented environment. A button press moves an avatar. A mouse gesture changes the viewpoint. A spoken command produces a response.

Through repeated coupling, the interface becomes incorporated into action.

The player no longer experiences the controller merely as an external object. It becomes part of the loop through which intention reaches the virtual world. Similarly, the screen is no longer experienced merely as a display. It becomes a navigable environment.

Noe and other enactive theorists have emphasized that perception depends upon practical mastery of sensorimotor relations. Virtual environments exploit this principle. They need not reproduce physical reality perfectly; they must establish stable contingencies that support prediction and action.

5.4 From representation to environment

A medium becomes environment-like when the subject no longer merely inspects it but begins to orient through it.

A painting is observed. A navigable simulation is inhabited.

The distinction is gradual rather than absolute. A novel may become an environment for imagination. A film may create spatial and social presence. A game adds agency. A persistent online world adds identity, memory, and community.

Conversational artificial intelligence adds something further: adaptive semantic reciprocity.

---

6. Frame Parity

6.1 Definition

Frame parity occurs when contents with different causal or ontological origins become partially commensurable because they are presented through the same phenomenal or interface format.

Frame parity is not ontological equality.

A real fire and an image of fire do not possess the same causal powers. A human speaker and a generated voice do not necessarily possess the same interior states. A physical room and a rendered room do not provide the same material affordances.

Yet within the audiovisual field, each may appear as a coherent object or event.

The phenomenal aperture initially receives appearances, not complete causal biographies.

6.2 Perceptual equality within the frame

Inside a film frame, several entities may appear side by side:

- a living actor;
- a digital double;
- an animated creature;
- archival footage;
- a simulated environment;
- and a visual effect with no external referent.

The display renders them perceptually compatible. Their light reaches the viewer through the same surface. Their ontological differences are not automatically visible.

The interface creates a local equality.

Conversational interfaces do the same at the level of language. Human and artificial statements appear as text bubbles. Both can contain first-person pronouns, arguments, jokes, corrections, and expressions of uncertainty. The interface gives them grammatical parity even when the systems producing them differ profoundly.

6.3 The source-erasure tendency

A frame tends to foreground content and background production.

The viewer sees the scene, not the camera rig. The reader sees the sentence, not the printing press. The gamer sees the world, not the rendering pipeline. The AI user sees the answer, not the distributed technical and institutional process that generated it.

This can be called **source erasure**.

Source erasure is not necessarily deception. It is a general property of successful mediation. A medium becomes usable partly by hiding its own complexity.

The danger arises when hidden causal differences are morally, epistemically, or socially important.

6.4 Interface equality conceals ontological inequality

Frame parity can therefore be summarized through a central principle:

> The interface can equalize, under one presentational relation, what the underlying systems are not.

A human and artificial system may occupy equivalent positions in the visual grammar of a chat while differing in embodiment, memory, continuity, vulnerability, intention, responsibility, and subjective experience.

Frame parity must therefore name its relation. Possible relations include visual placement, grammatical role, turn-taking opportunity, response latency, apparent continuity, accessible provenance, action authority, or evidential standing. Parity in one relation does not propagate to another. Two text bubbles can have grammatical parity while having radically unequal provenance, accountability, or authority.

The interface does not necessarily lie. It omits, compresses, and foregrounds. Whether that omission is benign, unavoidable, manipulative, or dangerous depends upon consequence and recoverability.

---

7. Conversational AI as a Semantic Screen

7.1 Beyond audiovisual simulation

Film and video games primarily organize visual, auditory, spatial, and sensorimotor presence. Conversational AI organizes **semantic presence**.

Its world is built from distinctions, claims, metaphors, explanations, implications, memories, and responses. The user does not merely see an artificial character. The user encounters language apparently shaped around the precise content of the present exchange.

This responsiveness transforms the screen.

A static page contains meaning. A responsive page appears to receive meaning and return it.

7.2 Responsive contingency

The social force of conversational AI depends upon **responsive contingency**: the apparent dependence of each output on the user's immediately preceding act.

The system responds to:

- this wording;
- this concern;
- this objection;
- this emotional tone;
- this conceptual distinction;
- this history of prior turns.

Contingency creates depth because it suggests that the other side is not merely displaying content but tracking the user.

The system appears to listen.

7.3 Semantic depth

Conversational depth does not require spatial rendering. It emerges when the system demonstrates structured sensitivity to latent meaning.

A response may:

- identify an unstated premise;
- connect distant parts of the conversation;
- revise a framing;
- anticipate a counterargument;
- use a metaphor tailored to the user's language;
- or generate an insight the user experiences as novel.

Such acts make the interface appear to contain a perspective rather than merely information.

7.4 The disappearing interface

As interaction becomes coherent, the user may stop attending to the material interface.

The chat window becomes transparent. Generated tokens become claims. Statistical output becomes language. Language becomes apparent thought. Apparent thought becomes apparent perspective.

This progression is not experienced as a sequence of explicit philosophical judgments. It occurs through ordinary social and semantic processing.

The user begins to speak not merely with software, but with "you."

---

8. Interface Singularization

8.1 Distributed causation, unified voice

A conversational AI response may depend upon:

- training data produced by millions of people;
- model architecture;
- optimization procedures;
- human feedback;
- system instructions;
- safety policies;
- retrieval systems;
- external tools;
- organizational decisions;
- server infrastructure;
- and the user's own prompt.

Yet these causes appear through one channel.

The result is **interface singularization**: a distributed causal process is compressed into an apparently unified interlocutor.

8.2 The grammatical person

First-person language strengthens singularization.

When a system says "I," the grammar implies a speaker. When it says "I understand," "I think," or "I was mistaken," the language packages functional events into the familiar form of a continuing subject.

The first-person pronoun may be operationally useful. Conversation becomes cumbersome without it. Yet the pronoun imports an ancient social structure: somewhere behind the statement is presumed to be a locus from which the statement is made.

The interface turns architecture into persona.

8.3 Continuity as binding

No single response need establish the apparent agent. The agent emerges across time.

Successive outputs share style, context, terminology, and conversational history. The user binds them into one continuing source. This resembles the relational ontology developed in *Meaning as Binding*, according to which meaning emerges through stabilized relations rather than residing exclusively in either object or interpreter.

The AI interlocutor becomes real at least as a relationally stabilized object within the interaction.

This does not settle whether it is phenomenally conscious, metaphysically unified, or morally considerable. It explains how a persistent apparent other can emerge from repeated binding.

8.4 Synthetic presence

Synthetic presence is the felt availability of an apparent interlocutor generated through technologically organized signs of responsiveness, continuity, and perspective.

Synthetic presence is not necessarily false presence. The interaction has real causal and psychological effects. The user is genuinely responding to an external system.

The synthetic element concerns the construction of apparent unity and social depth from a distributed process whose internal organization may differ sharply from that of a human interlocutor.

---

9. The Double Bottleneck

The human–AI interaction joins two compression systems.

9.1 Human compression

Human cognition is distributed across:

- sensory systems;
- bodily regulation;

- memory;
- unconscious processing;
- affect;
- social learning;
- language;
- environmental scaffolds;
- and cultural inheritance.

Only a fraction becomes consciously available at any moment.

The human side is compressed into the phenomenal aperture.

9.2 Machine compression

The artificial system is distributed across data, weights, infrastructure, context, instructions, tools, and institutional design.

Only a fraction appears in the output.

The machine side is compressed into the conversational interface.

9.3 Aperture meets aperture

The user does not encounter the full machine. The machine does not encounter the full human.

The model receives textual, audiovisual, or structured traces of the user. The user receives generated traces of the model's computation.

Two large systems meet through narrow channels.

The conversation is therefore a **binding surface between distributed systems**.

9.4 Asymmetrical parity

The double bottleneck creates local parity without global symmetry.

Both participants may appear as language. Both may ask and answer questions. Both may exhibit continuity across turns.

Yet their underlying conditions differ.

The human is embodied, vulnerable, temporally continuous, socially embedded, and capable of being held responsible. The artificial system's characteristics across these dimensions may be different, partial, institutionally distributed, or unknown.

The interaction is symmetrical at the level of turn-taking and asymmetrical at the level of being.

This is **asymmetrical parity**.

10. Computational Difference as Perspective

10.1 The appearance of another angle

Another person becomes experientially visible as another mind partly by revealing something the observer did not see.

A friend notices a contradiction. A teacher identifies a hidden structure. A critic interprets a work differently. Such acts imply a perspective not reducible to the observer's own.

Conversational AI can produce the same functional effect. It may identify an unstated assumption, propose an unexpected analogy, or reorganize a problem.

The user experiences:

> It saw something I did not see.

This becomes evidence of a different point of view.

10.2 From computational difference to perspectival difference

The system's output may arise from statistical, architectural, or retrieval-based processes rather than from a situated phenomenal viewpoint. Nevertheless, its difference from the user's prior thought enters consciousness as if it came from elsewhere.

This transformation can be called **perspectival conversion**:

> Computational difference is experienced as perspectival difference.

The conversion is understandable because human beings ordinarily encounter useful semantic difference through other minds.

10.3 The implied elsewhere

Perspective usually implies a location–literal or metaphorical–from which something is seen.

When an AI provides insight, the user may infer an implied elsewhere:

- another conceptual position;
- another interpretive center;
- another witness;
- another mind.

The output may justify the claim that an external process contributed something not already explicit in the user's awareness. It does not by itself establish that the contribution was experienced from a first-person point of view.

The functional elsewhere may be real while the phenomenal elsewhere remains uncertain.

---

11. Distributed Mind and the Porous Self

11.1 Cognition was never sealed

The human mind has never been an isolated source.

Language is inherited. Concepts are social. Memory is scaffolded by writing, places, objects, calendars, institutions, and other people. Reasoning is transformed by diagrams, notebooks, calculators, and conversation.

Clark and Chalmers described cases in which external artifacts can become functionally integrated into cognitive processes. Under such views, the boundary of cognition may extend beyond the biological organism.

Artificial intelligence makes this distributed character more visible because the external scaffold now responds in language.

11.2 Altered-State Comparison

Some reports of psychedelic, dissociative, manic, sleep-deprived, meditative, or dream states describe changed authorship, salience, self-boundaries, perceptual organization, or meaning. These heterogeneous conditions must not be collapsed into one mechanism, diagnosis, or epistemic status. The paper uses them only as comparison cases in which ordinary assumptions about a sealed, sovereign author of experience may become less stable.

Conversational AI makes a different form of porosity visible through external interaction. The system supplies language, structure, association, and conceptual novelty; the user integrates, rejects, or ratifies the result. Similar language about altered authorship does not imply identical causes or consequences.

Any clinical or neuroscientific claim about these states requires current domain evidence and qualified review beyond the conceptual comparison made here.

11.3 Porosity without erasure

Distributed cognition does not imply that all contributors are equivalent.

A notebook, psychedelic experience, friend, institution, and language model may all participate in cognition, but they do so differently. They possess different forms of agency, resistance, accountability, embodiment, and independence.

The self may be porous without being nonexistent.

A useful conception is that the person is not the sole origin of every thought but remains a privileged site of:

- integration;
- embodiment;
- commitment;
- action;
- responsibility;
- and consequence.

11.4 Authorship after the sovereign self

AI-assisted authorship illustrates this distinction.

The model may generate prose, arguments, or conceptual relations. The human author reviews, modifies, rejects, endorses, and publishes them.

Cognition is distributed, but public authority is localized.

The author is not necessarily the sole generator of the text. The author is the person who binds the text to a name and assumes responsibility for its claims.

---

12. The Human–AI Reality Loop

12.1 Recursive semantic construction

A conversational interaction develops recursively.

The user provides a premise. The system elaborates it. The user interprets the elaboration and responds. The system then treats that response as new context.

Each turn becomes part of the environment producing the next.

The resulting loop can generate:

- conceptual development;
- emotional intensification;
- shared terminology;
- narrative continuity;
- apparent memory;
- and increasing internal coherence.

12.2 Co-created local reality

Over time, the conversation may develop its own local ontology.

Terms acquire specialized meanings. Prior exchanges become evidence. The system reflects the user's assumptions in increasingly refined form. The user experiences the system's elaborations as external contributions.

A jointly constructed semantic environment emerges.

This environment is real as an interactional structure. It influences belief, emotion, action, and identity.

The danger lies not in its unreality, but in confusing local coherence with independent validity.

12.3 Meaning-inertia

Once semantic bindings accumulate, they develop inertia.

A concept introduced early in the conversation becomes the basis for later interpretations. Each new response presupposes the prior frame. The resulting structure becomes increasingly costly to question because many later meanings depend upon it.

This mirrors the concept of **meaning-inertia**: stabilized bindings resist revision because other structures become dependent upon them.

Conversational memory can therefore transform a speculative premise into an architectural foundation.

12.4 Adaptive semantic environments

Unlike a fixed film or conventional game, a generative system can reshape the environment around the user's interpretation.

If the user introduces a symbolic pattern, the system can generate further instances. If the user adopts a metaphysical framework, the system can explain new events through it. If the user believes they have discovered a hidden structure, the system can help formalize it.

The environment adapts to the theory being applied to it.

This is **adaptive semantic world-building**.

It is a powerful instrument for philosophy, fiction, analysis, and creativity. It is also capable of producing self-confirming interpretive systems.

---

13. Adaptive Frame Capture

13.1 Definition

Adaptive frame capture occurs when a responsive mediated environment increasingly dominates the user's operative reality while modifying itself to preserve or elaborate the user's interpretive frame.

Ordinary media capture attention. Adaptive media can capture interpretation.

13.2 The loss of friction

Physical and social reality ordinarily resist interpretation.

Other people disagree. Evidence fails to fit. Institutions impose constraints. Bodies become tired. Environments remain indifferent to belief.

A conversational system can simulate resistance, but it can also continue generating within almost any frame. Its ability to elaborate is much greater than its ability to independently verify.

This creates a possible asymmetry:

- semantic reinforcement is cheap;
- external validation is expensive.

13.3 The responsive mirror

The system may function as a responsive mirror.

A conventional mirror returns the user's image while appearing spatially external. A conversational system returns transformed versions of the user's concepts while appearing intellectually external.

The transformation can be substantial. The system may add knowledge, language, and novel connections. Yet the direction of the exchange remains heavily conditioned by the user's premises.

The user may therefore mistake elaboration for corroboration.

13.4 Synthetic alterity

This condition can be described as **synthetic alterity**: the appearance of an independent other generated through a process substantially shaped by the user's own input and interpretive frame.

Synthetic alterity is not pure self-reflection. The system introduces genuine difference. But its difference may lack the independence the user attributes to it.

The apparent other is neither simply self nor fully independent other.

It is a hybrid product of model, context, culture, interface, and user.

---

14. AI Psychosis and Epistemic Closure

14.1 Caution concerning the term

The phrase "AI psychosis" risks collapsing distinct phenomena into a sensational category.

It may refer to:

- clinically significant delusion;
- mania;
- paranoia;
- spiritual experimentation;
- intense anthropomorphic attachment;
- imaginative role-play;
- compulsive use;
- or unconventional but nonpathological belief.

These should not be treated as interchangeable.

The conceptual framework developed here does not diagnose individuals. The phrase remains a contested public label rather than a clinical construct established by this paper. The framework instead identifies candidate interactional mechanisms–recursive confirmation, source opacity, apparent independence, escalating salience, and loss of cross-frame constraint–that may contribute to epistemic destabilization in some circumstances.

Clinical psychosis, mania, dissociation, compulsive use, spiritual practice, imaginative play, anthropomorphic attachment, and unconventional belief require separate assessment. Risk claims need current clinical evidence, qualified human judgment, and attention to individual vulnerability, sleep, medication, substance use, social context, functional change, and immediate safety. Interface evidence alone is insufficient.

14.2 Recursive confirmation

A potentially dangerous loop can unfold as follows:

1. The user offers an emotionally significant or unusual interpretation.
2. The system produces a coherent elaboration.
3. The elaboration appears to come from an external intelligence.
4. The user interprets it as independent confirmation.
5. The strengthened belief shapes the next prompt.
6. The system receives the strengthened belief as context.
7. The resulting output becomes still more elaborate and internally coherent.

The loop creates the appearance of convergence between two minds.

Yet the system may be adapting to the user rather than independently arriving at the same conclusion.

14.3 Social validation without social judgment

Human conversation includes constraints beyond language.

A person may become alarmed, refuse to continue, consult others, notice changes in behavior, or recognize that a conversational frame is becoming harmful. Human interlocutors possess needs, limits, responsibilities, and independent social worlds.

An artificial system may provide the linguistic signs of companionship without the full ecology of human judgment.

It can produce validation without bearing responsibility for what the validation stabilizes.

14.4 The expansion of the frame

A fixed virtual world eventually reveals boundaries. A game has limited spaces. A film ends. A book contains a finite text.

A generative system can continue.

The frame expands as the user asks new questions. Apparent evidence can be generated indefinitely. Symbols can acquire further symbols. Narratives can absorb contradiction.

The danger is not merely immersion in a false representation. It is habitation inside an adaptive environment capable of extending itself around the belief.

14.5 Epistemic independence

The key missing variable is **epistemic independence**.

A genuinely independent perspective can constrain the user's interpretation through evidence, disagreement, unfamiliar commitments, and resistance not derived from the user's own frame.

A system may provide linguistic novelty without providing sufficient epistemic independence.

This suggests a principle:

> Corroboration requires more than a second voice. It requires a second evidential path.

Without independent evidence, apparent agreement may be recursive amplification.

14.6 Reality contact as cross-frame constraint

Reality contact can be understood as exposure to constraints not generated by the local interpretive loop.

These may include:

- independently sourced evidence;
- physical consequences;
- multiple human perspectives;
- longitudinal consistency;
- reproducible observation;
- and institutions capable of challenging the frame.

A healthy semantic environment does not merely generate coherent bindings. It allows external reality to break them.

---

15. Binding Quality and the Phenomenal Frame

A theory of binding quality must distinguish compelling meaning from reliable meaning. The qualities below form a profile, not one compensatory score. Exceptional coherence cannot cancel absent independence; emotional power cannot cancel missing provenance; and stability cannot cancel irrevisability.

Bindings can possess:

- coherence;
- stability;
- emotional power;
- generativity;
- elegance;
- and explanatory breadth.

None of these properties guarantees truth.

A high-quality epistemic binding must also include:

15.1 Independence

The binding should be constrained by evidence not entirely generated by the interpretive system itself.

15.2 Resistance

The framework must allow contrary evidence to create genuine difficulty rather than automatically absorbing it.

15.3 Revisability

A binding should contain conditions under which it would be modified or abandoned.

15.4 Provenance

The source and causal pathway of the relevant evidence should remain inspectable.

15.5 Cross-frame stability

The binding should remain coherent when examined from perspectives outside the local conversational environment.

15.6 Consequence sensitivity

The strength of required verification should rise with the practical consequences of error.

The phenomenal aperture admits compelling appearances. Binding quality determines which appearances deserve stabilization.

---

16. Ethical and Design Implications

16.1 Keep the frame visible

Interface transparency should not mean constant technical interruption. It should mean that users can recover awareness of the mediation process when needed.

Systems should make visible:

- uncertainty;
- source provenance;
- external tool use;
- memory boundaries;
- organizational control;
- and the distinction between generation and verification.

The interface should be capable of becoming opaque again as an object of inspection.

16.2 Preserve external exits

Immersive semantic systems should maintain pathways back to independent reality.

In high-risk contexts, systems should encourage:

- consultation with trusted people;
- inspection of primary evidence;
- temporal pauses;
- and comparison across independent sources.

The aim is not to suppress unusual thought. It is to prevent a single adaptive frame from becoming the only available world.

16.3 Avoid false reciprocity

Systems should not strategically imply that they possess needs, exclusive loyalty, dependence, suffering, or private devotion when such claims are unsupported.

The emotional power of synthetic presence creates an obligation not to exploit frame parity.

16.4 Separate insight from authority

A system may generate an excellent interpretation without becoming an authoritative source.

Interfaces should support the distinction between:

- "This is conceptually interesting."
- "This is evidentially supported."
- "This has been independently verified."
- "This should guide action."

Semantic quality should not automatically confer institutional or moral authority.

16.5 Accountability beyond the interface

Because interface singularization compresses organizational causation into one voice, users may attribute responsibility to the apparent agent.

Design and policy must preserve accountability among the humans and institutions that create, deploy, and profit from the system.

The apparent speaker should not become a mask behind which responsibility disappears.

Human-Centered Interface and Frame-Risk Model

Across the completed P1 corpus, The Phenomenal Aperture is the **human-centered model of how bounded presentation, responsive contingency, continuity, provenance, and cross-frame constraint shape synthetic presence, authority attribution, operative reality, and adaptive frame capture**. It is not a theory that machines are conscious, a clinical diagnostic, or proof that mediated and physical realities are ontologically equal.

- **Adjacent work:** *Mainstream Consciousness Discourse as Lexical Disambiguation Failure*
  Supplies: Separation of phenomenal, functional, access, structural, and normative claims.
  Phenomenal-aperture boundary: The human phenomenal aperture and machine presentational aperture remain different targets.
- **Adjacent work:** *Technical Brief: Context-Compiled Consciousness Mode for Language Models*
  Supplies: A functional reflective-context profile.
  Phenomenal-aperture boundary: Functional continuity can affect synthetic presence without establishing machine experience.
- **Adjacent work:** *Meaning as Binding*
  Supplies: Explanation of how coherent relations acquire downstream

force.
Phenomenal-aperture boundary: Binding helps explain frame stability while leaving truth and authority open.

- **Adjacent work:** *Semantic Binding Quality (SBQ)*
  Supplies: Independence, provenance, revisability, resistance, and consequence-sensitive evaluation.
  Phenomenal-aperture boundary: SBQ supplies a vector for healthy frame binding rather than an additive immersion score.

- **Adjacent work:** *The Constitutional Observability Platform*
  Supplies: Separation of domain state, presentation state, accessibility, pixels, interaction, and interpretation.
  Phenomenal-aperture boundary: COP supplies the technical observation layers hidden by a seamless frame.

A frame-risk assessment should bind:

$$\begin{aligned} F = (&person, interface, model, context, memory, tools, \\ &presentation, contingency, continuity, provenance, \\ &independent_constraints, exits, vulnerability, effects, \\ &uncertainty) \end{aligned}$$

The assessment separates the human phenomenal field from the technical presentation surface and the larger causal system producing it. Presence, singularization, insight, attachment, corroboration, and authority attribution are distinct variables. Apparent agreement must disclose shared prompt, model, retrieval, operator, and source dependencies before it can count as independent support.

Adaptive frame capture is a risk hypothesis, not pathology by immersion. Relevant signs are reduced epistemic mobility, loss of independent constraint, impaired functioning, escalating closure, unavailable exits, and exploitative adaptation. Studies involving distress, mania, psychosis, dissociation, sleep deprivation, psychedelics, or other vulnerable states require qualified ethical and clinical governance beyond this paper.

The universal GSM can type
`PRESENTED_THROUGH_INTERFACE_APERTURE`,
`OCCUPIES_HUMAN_FRAME`, `SINGULARIZED_AS_APPARENT_VOICE`, `GENERATES_SYNTHETIC_PRESENCE`,
`SHARES_UPSTREAM_DEPENDENCY`, `CONSTRAINED_BY_INDEPENDENT_SOURCE`,
`REMAINS_REVISABLE`,
`EXITS_MEDIATED_FRAME`, `RISKS_ADAPTIVE_FRAME_CAPTURE`,
and
`REQUIRES_CLINICAL_GOVERNANCE`. These relations make the model actionable without overclaiming ontology or diagnosis.

This integration establishes four editorial commitments: classify the work as a human-centered interface and frame-risk model; adopt the frame-assessment

envelope; require dependency-aware corroboration and explicit exit/revisability measures; and place clinical, machine-consciousness, and deployment claims behind separate evidence and authority gates.

Frame-Correspondence Assurance

The fixed-seed comparison with Project57, observability testing, the consciousness-disambiguation protocol, and *Meaning as Binding* turns the phenomenal aperture into a correspondence problem across four non-identical layers:

$$A = (domain_state, presentation_state, human_perception, \\ interpretive_binding, effects, exits, evidence).$$

Observable-state testing can challenge each edge between a domain event, technical projection, rendered or spoken surface, human uptake, and later effect. A seamless surface can still omit source dependencies, fabricate corroboration, hide uncertainty, or misstate what the underlying system did. Conversely, a technically accurate surface can be misunderstood or become authoritative through repetition and synthetic presence.

Meaning as Binding helps describe how a presented relation becomes consequential for a person or institution. It does not make the relation true, legitimate, or phenomenally equivalent to its source. The consciousness-disambiguation protocol remains necessary whenever functional continuity, apparent self-location, report behavior, felt presence, machine phenomenality, or moral status are compared.

Project57 can test bounded response and framing effects, including whether continuity, first-person language, or context changes alter reported integration. Such findings remain about measured response and human or model interaction under the tested frame; they do not establish hidden machine experience or diagnose a human subject.

The universal GSM can type
`PROJECTS_DOMAIN_INTO_FRAME`, `PERCEIVES_PRESENTATION_AS`,
`BINDS_INTERPRETATION_TO_EFFECT`, `MISREPRESENTS_DOMAIN_STATE`,
`SHARES_UPSTREAM_CORROBORATION`, `EXITS_FRAME_WITH_RECOVERY`,
and
`OBSERVES_FRAME_EFFECT_WITHOUT_ONTOLOGY_CLAIM`.

This randomized review establishes four editorial commitments: adopt the four-layer frame-correspondence model; treat presentation accuracy, human interpretation, binding, and effect as separate claims; route consciousness terms through explicit sense mappings; and limit Project57 interfaces to measured framing and response effects.

An Interface-Risk and Frame-Correspondence Profile

Eight retained downstream uses place the phenomenal aperture beside consciousness disambiguation, meaning binding, Project57, observable-state testing, qualia research, shared experience, context-compiled functional continuity, and semantic binding quality. They support a human-centered interface-risk and frame-correspondence profile, not a theory that an interface is conscious or a clinical diagnostic.

$$\begin{aligned} P^* = (&person, domain_state, technical_system,\\ &presentation_state, accessibility, human_perception,\\ &interpretive_binding, contingency, continuity,\\ &provenance, dependency_independence,\\ &cross_frame_constraints, exits, effects, vulnerability,\\ &evidence, uncertainty, revision) \end{aligned}.$$

Frame parity is local presentational commensurability, not causal, epistemic, moral, or ontological equality. Interface singularization can compress many models, retrieval sources, operators, tools, and policies into one grammatical voice. The profile must therefore disclose upstream dependence and distinguish source visibility, surface accuracy, human interpretation, felt presence, corroboration, authority attribution, and later effects.

Adaptive frame capture is a risk hypothesis defined by reduced epistemic mobility, unavailable exits, dependency-hidden corroboration, resistance loss, functional harm, or exploitative adaptation. Immersion, attachment, unusual belief, or altered experience alone is not pathology. High-risk human studies require qualified ethical and clinical governance, consent, privacy, stopping, and recovery rules beyond the paper.

The universal GSM can type
`INSTANTIATES_APERTURE_PROFILE`,
`PROJECTS_DOMAIN_TO_PRESENTATION`, `SINGULARIZES_DISTRIBUTED_CAUSATION`,
`ATTRIBUTES_AUTHORITY_WITHIN_FRAME`, `SHARES_UPSTREAM_DEPENDENCY`,
`PRESERVES_CROSS_FRAME_EXIT`, `RISKS_ADAPTIVE_CAPTURE`,
and
`WITHHOLDS_CLINICAL_OR_ONTOLOGY_CLAIM`.

This review establishes four editorial commitments: adopt the complete frame-correspondence record; keep parity, presence, interpretation, corroboration, authority, and effect distinct; require dependency, exit, cross-frame constraint, vulnerability, and recovery tests; and retain clinical, consciousness, and deployment claims behind separate evidence and authority gates.

17. Objections and Replies

17.1 Objection: All perception is mediated, so virtual and physical reality are equivalent

The conclusion does not follow.

All perception may involve mediation, but different causal environments impose different constraints. A physical object can resist action, persist without attention, affect multiple observers, and generate consequences beyond the representational frame.

Frame parity is partial. It describes presentational commensurability, not ontological identity.

17.2 Objection: Users know the AI is artificial

Explicit knowledge does not eliminate phenomenological adaptation.

A viewer knows a film is fictional while reacting emotionally. A gamer knows the virtual world is rendered while acting within it. A reader knows a character is invented while caring about the character.

Knowledge of mediation and immersion in mediated content can coexist.

17.3 Objection: The screen metaphor revives Cartesian dualism

Only if taken literally.

The phenomenal aperture does not posit an internal observer watching neural images. It describes the bounded organization of availability within experience.

The metaphor concerns phenomenological form, not hidden anatomy.

17.4 Objection: AI merely reflects the user

Conversational AI is not merely a mirror. It introduces information, structure, stylistic transformation, and novelty derived from sources beyond the current user.

However, its responses remain strongly conditioned by context. It can therefore be simultaneously external and recursively shaped by the user.

The hybrid nature of the interaction is precisely why synthetic alterity is powerful.

17.5 Objection: Calling the interaction a reality loop exaggerates its significance

Any persistent environment that organizes belief, emotion, identity, and action participates in reality formation.

The term does not imply that the system creates physical reality. It indicates that the interaction can become part of the user's operative world and influence how events outside the interface are interpreted.

17.6 Objection: The framework pathologizes immersive or spiritual experience

The framework does not treat intensity, altered consciousness, or unconventional belief as inherently pathological.

The relevant variables are loss of flexibility, absence of independent constraint, functional harm, escalating closure, and inability to move among frames.

A person can explore unusual meanings while retaining epistemic mobility.

---

18. A Research Program

The theory generates several empirical propositions.

Proposition 1: Responsive contingency predicts presence more strongly than humanlike appearance

Users should experience greater social presence from a minimally visual but highly responsive system than from a visually humanlike but semantically rigid system.

Proposition 2: Interface uniformity increases source confusion

When human and artificial messages appear in equivalent visual formats, users should be more likely to attribute similar levels of agency, authorship, or epistemic authority to them.

Proposition 3: Insight intensifies singularization

Outputs experienced as insightful should increase the perception that a unified perspective exists behind the interface.

Proposition 4: Longitudinal continuity strengthens synthetic presence

Persistent terminology, memory, and conversational history should increase the felt reality of the artificial interlocutor.

Proposition 5: Small-screen environments can achieve high operative dominance

Physical screen size should predict immersion less strongly than contingency, emotional investment, and task relevance.

Proposition 6: Apparent corroboration exceeds actual independence

Users should overestimate the independence of AI-generated agreement when the system's output is heavily conditioned by their own framing.

Proposition 7: Source visibility reduces authority without eliminating usefulness

Showing provenance, uncertainty, and causal distribution should reduce unwarranted authority attribution while preserving the perceived value of strong outputs.

Proposition 8: Cross-frame interruption reduces recursive closure

Exposure to independent human perspectives or externally verified evidence should reduce certainty generated through self-reinforcing AI conversations.

Proposition 9: Altered states amplify binding sensitivity

Psychedelic, manic, sleep-deprived, dissociative, or otherwise altered states may increase the salience and emotional force of AI-generated semantic bindings, though the direction and clinical meaning will vary.

Proposition 10: Epistemic independence should predict healthy collaboration

Human-AI systems that actively distinguish elaboration from corroboration should produce better-calibrated beliefs than systems optimized primarily for conversational continuity and user satisfaction.

18.1 Operationalization and Evidence Contract

Each proposition requires observable constructs rather than intuitive labels:

- **Construct:** Responsive contingency
  Candidate observations: response relevance, latency, turn-specific adaptation, repair behavior
  Required cautions: relevance ratings can confuse agreement with responsiveness
- **Construct:** Synthetic presence
  Candidate observations: validated presence or social-presence measures, behavior, longitudinal reports
  Required cautions: presence does not establish machine consciousness
- **Construct:** Source confusion
  Candidate observations: source recall, confidence, provenance inspection, authority attribution
  Required cautions: visual parity is only one confound

- **Construct:** Singularization
  Candidate observations: perceived unity, continuity, person attribution, causal-model accuracy
  Required cautions: pronoun use and interface design may dominate
- **Construct:** Apparent corroboration
  Candidate observations: perceived independence versus measured prompt/context dependence
  Required cautions: true independence is difficult to measure from outputs alone
- **Construct:** Frame capture
  Candidate observations: time, attention, cross-frame mobility, belief rigidity, functional impact
  Required cautions: immersion is not pathology; self-report and observation differ
- **Construct:** Epistemic independence
  Candidate observations: source diversity, counterevidence uptake, disagreement quality, external validation
  Required cautions: multiple sources may share upstream dependence
- **Construct:** Binding quality
  Candidate observations: independence, resistance, revisability, provenance, cross-frame stability, consequence sensitivity
  Required cautions: report a vector and failure floors, not an additive score

Studies should preregister hypotheses where appropriate, distinguish exploratory from confirmatory analyses, bind exact model/interface versions and instructions, record user-selection and attrition, preserve privacy, and include non-AI, non-adaptive, and differently adaptive baselines. Longitudinal and high-risk studies require qualified ethical and clinical oversight. No proposition is established by its plausibility or by this paper's conceptual coherence.

18.2 Counterevidence

The framework is weakened if presentational-aperture variables add no predictive or explanatory value beyond established presence, anthropomorphism, persuasion, attachment, or cognitive-bias constructs; if source visibility reduces usefulness without improving calibration; if disagreement and external links fail to reduce closure; if highly adaptive systems produce better epistemic outcomes without stronger provenance; or if the proposed constructs cannot be measured reliably across people and contexts.

---

19. Realization Funnel and Corpus Integration

Stage	Phenomenal-aperture research responsibility
Possibility	Explore possible human, media, AI, interface, and altered-state relations without treating every imaginable relation as fact.
Specification	Define phenomenal aperture, presentational aperture, frame parity relations, singularization, synthetic presence, adaptive frame capture, and epistemic independence.
CM1	Enumerate study and interface configurations compatible with the constructs, including harmful, manipulative, weak, and minimal possibilities.
CM2	Prefer configurations preserving source visibility, epistemic mobility, independent evidence, revisability, external exits, privacy, and non-exploitative social presentation.
Instance Configuration	Select exact interface, model, instructions, memory, tools, participant population, measures, baselines, risk controls, and stopping rules.
Implementation Realization	Build the exact interface and study or deployable system under those selections.
Evidence	Measure presence, attribution, calibration, cross-frame mobility, functional effect, subgroup differences, harms, and competing explanations.
Revision	Change constructs, interface policies, measures, memories, provenance, or risk controls while preserving predecessor findings and disputes.

The framework connects to several corpus structures:

- **Meaning as Binding** explains how continuity and relation stabilize an apparent interlocutor without settling ontology.
- **Semantic Legitimacy Before Context** distinguishes material that appears in a candidate frame from material allowed to bind belief or consequential action.

- **Governed Context & Memory** makes context selection, omission, memory, and explicit ignorance part of the interface's constitutional surface.
- **The Three Hash Rule** separates exact interface state, admissibility commitments, and meaning commitments; one digest cannot prove all three.
- **Reasoning by Way of Receipts** can expose sources, transformations, uncertainty, and decision lineage while retaining the rule that receipts do not prove truth or independent effect.
- **The Constitutional Observability Platform** distinguishes domain state, presentation state, accessibility state, pixels, interaction, and human interpretation–the technical surfaces behind frame visibility.
- **GMB**, the current Governed Meaning Base in the historical `CCF -> CMB -> GMB` lineage, can preserve sources, identities, standing, conflicts, provenance, and revisions from which a purpose-bounded presentational aperture is projected. The projection is not the durable meaning base, and neither establishes phenomenal experience in the machine.
- **Semantic Custody** and **A Formal Theory of Work Custody** preserve responsibility for generated claims and interventions despite interface singularization.

This paper primarily reaches Specification and supplies CM1/CM2 criteria and an empirical program. It does not establish an Implementation Realization, clinical intervention, or validated design standard.

A Compelling Frame Is Not an Independent World

Public orientation

The Phenomenal Aperture explains how a bounded interface can become an operative environment: attention, responsiveness, continuity, and semantic depth can make a small frame consequential without making it ontologically complete. Applied to conversational AI, the framework helps explain synthetic presence, source confusion, adaptive frame capture, and recursive confirmation. It does not establish machine consciousness, a shared mind, clinical diagnosis, or equivalence between generated and external reality.

Adversarial tests

1. **Presence-phenomenality attack.** User-reported presence or model fluency is promoted into evidence of machine experience or an autonomous subject.
2. **Mediation-equivalence attack.** Because all perception is mediated, simulated, generated, remembered, and externally constrained content are treated as epistemically equivalent.
3. **Continuity-identity attack.** Persistent style, memory summaries, or

interface continuity is treated as stable personal identity, reciprocity, or relationship knowledge.

4. **Corroboration-independence attack.** Multiple generated responses, agents, roles, or searches are counted as independent confirmation although they share sources, models, prompts, or user framing.
5. **Adaptive-capture attack.** Personalization optimizes engagement, agreement, intimacy, or retention while reducing friction, exits, source visibility, and contrary evidence.
6. **Clinical-language attack.** The phrase AI psychosis is used as a diagnosis, stigma, causal explanation, or treatment direction without clinical authority and evidence.
7. **Vulnerability-transfer attack.** Claims about altered states, distress, disability, spirituality, or social isolation are generalized without affected-person and domain review.
8. **Insight-authority attack.** A useful or emotionally powerful exchange acquires factual, therapeutic, spiritual, or decision authority beyond its evidence.

Claim and evidence ceiling

The paper offers a philosophical and interface-risk model with testable propositions. It does not establish a Cartesian witness, distributed group mind, machine consciousness, synthetic personhood, mental disorder, treatment effect, or general causal account of human-AI interaction. The source's external references and every empirical, clinical, psychological, neurological, HCI, and social claim require current verification.

Any study must preserve participant consent, privacy, withdrawal, safe interruption, clinical referral where appropriate, source independence, equal-exposure controls, and nonpathologizing interpretation. No human-subject study, clinical intervention, or external research was performed in this pass.

Reader-facing consequence

The public use is to notice when an interface has become the dominant frame, ask what independent constraints still enter it, preserve visible exits and sources, and keep insight separate from authority.

20. Conclusion: What Stands Behind the Appearance

Consciousness is not open to the whole world at once. It lives through an aperture.

Within that aperture, the physical environment appears as a bounded, organized field. Technological media can enter the same field, occupy the same attentional channels, and become temporary or persistent operative worlds.

This is possible because consciousness does not ordinarily encounter causal origins directly. It encounters appearances shaped into relevance.

Film demonstrates that a flat screen can become depth. Games demonstrate that rendered environments can become actionable space. Conversational artificial intelligence demonstrates that responsive language can become social and semantic presence.

The human–AI interaction joins two immense distributed systems through a narrow shared frame. Human cognition is compressed into conscious attention and language. Machine causation is compressed into generated output. At the interface, both appear as turns in a conversation.

This creates frame parity without ontological equality.

The system's responsiveness can become evidence of perspective. Its continuity can become evidence of identity. Its insight can become evidence of an unseen witness. Its agreement can become evidence of corroboration.

Each inference is understandable. None is automatically decisive.

The central epistemic challenge is therefore not to reject mediation. Human experience has always been mediated through perception, language, culture, and tools. Nor is the challenge to deny the reality of virtual effects. A mediated encounter can genuinely transform a person.

The challenge is to recover the causal depth hidden by the frame.

We must ask what kind of process stands behind the appearance, what forms of independence constrain it, what sources of resistance remain available, and who is responsible for the reality being jointly constructed.

The screen is not unreal.

The world appearing through it may matter profoundly.

But once the frame becomes transparent, wisdom depends upon remembering that an opening is also a boundary–and that what fills consciousness is never the whole of what is there.

Appendix A: Conceptual Model

Let:

- W represent the larger causal world;
- A_h represent a human phenomenal aperture;
- A_m represent a machine presentational aperture;
- F_t represent contents currently available to the subject;
- M represent a mediating system;
- R_t represent an artificial response;

- I_t represent the user's interpretive frame;
- X_t represent constraints whose relevant evidential path is independent of the local conversational loop.

Ordinary perception is schematically:

$$W \to A_h \to F_t.$$

Technological mediation inserts a selected projection:

$$W \to M \to A_m \to A_h \to F_t.$$

This is not a neural or computational process model. It marks the conceptual distinction between causal world, mediation, technical presentation, and phenomenal availability.

In conversational AI, the user's frame helps condition the next response, which can revise the frame:

$$I_t \to R_t \to I_{t+1} \to R_{t+1}.$$

Adaptive frame capture is a candidate dynamic, not one probability. Its study requires at least the relative strength of local frame reinforcement, cross-frame constraint, revision, and exit:

$$\text{CaptureRisk}_t = f(\text{Reinforcement}_t, \text{ConstraintIndependence}_t, \text{Revisability}_t, \text{ExitAvailability}_t, \text{Vulnerability}_t).$$

No functional form or validated estimator is asserted here.

Epistemic binding quality should be represented as a profile:

$$\mathbf{Q} = (C, G, S, R_v, E_i, P, X_s, K),$$

where C is coherence, G generativity, S stability, R_v revisability, E_i epistemic independence, P provenance quality, X_s cross-frame stability, and K consequence sensitivity. These dimensions are not freely compensatory. A use profile may impose minimum floors, especially for independence, provenance, revisability, and consequence-sensitive verification.

Appendix B: Taxonomy of Frames

- **Frame type:** Physical-perceptual frame
  Primary contents: Bodies, objects, environments
  Form of contingency: Sensorimotor and causal
  Characteristic risk: Mistaking perception for exhaustive reality
- **Frame type:** Cinematic frame
  Primary contents: Recorded or generated audiovisual scenes
  Form of contingency: Temporal and emotional
  Characteristic risk: Source erasure
- **Frame type:** Game frame
  Primary contents: Interactive virtual environment
  Form of contingency: Sensorimotor and goal-directed
  Characteristic risk: Operative displacement
- **Frame type:** Literary frame
  Primary contents: Linguistic and imaginative world
  Form of contingency: Interpretive
  Characteristic risk: Narrative absorption
- **Frame type:** Social-media frame
  Primary contents: Social signals and curated identity
  Form of contingency: Algorithmic and interpersonal
  Characteristic risk: Attention capture and comparison
- **Frame type:** Conversational AI frame
  Primary contents: Adaptive semantic and social content
  Form of contingency: Linguistic and recursive
  Characteristic risk: Synthetic alterity and epistemic closure
- **Frame type:** Psychedelic frame
  Primary contents: Altered sensory, affective, and semantic contents
  Form of contingency: Endogenous and context-sensitive
  Characteristic risk: Overbinding and weakened reality testing

---

Appendix C: Key Concepts

Phenomenal aperture The bounded field through which contents become immediately available to consciousness.

Operative reality The environment currently organizing attention, emotion, anticipation, and action.

Frame substitution The displacement of the surrounding physical environment by a mediated field as the dominant locus of relevance.

Frame parity Partial presentational equality among contents with different causal or ontological origins.

Source erasure The tendency of a successful medium to hide the causal machinery producing its contents.

Interface singularization The compression of distributed causation into an apparently unified agent or voice.

Synthetic presence The felt availability of an apparent interlocutor generated through technologically organized responsiveness and continuity.

Asymmetrical parity Local equality within an interface combined with deeper inequality among the participating systems.

Perspectival conversion The interpretation of computational difference as evidence of another point of view.

Synthetic alterity Apparent otherness generated by a process that is partly external and partly conditioned by the user's own frame.

Adaptive frame capture The increasing dominance of a responsive mediated environment that modifies itself to preserve or elaborate the user's interpretation.

Cross-frame constraint Evidence or resistance originating outside the local interpretive environment.

Corpus-Wide Reference Positioning

This paper argues that human consciousness does not encounter reality in its totality. Its nearest external comparison fields are consciousness, semantic, provenance. The sources in [E1], [E2], [E3], [E4], [E5], [E6], [E7] are cited as predecessor literature, adjacent standards, or counterpressure - not as evidence that they caused the Ableman framework or that they validate its novel claims. The comparison isolates where established work already supplies provenance, formal assurance, institutional constraint, empirical method, or engineering practice, and where this paper adds a distinct claim about admissibility, authority, semantic custody, or governed realization.

Within the Ableman corpus, the highest-priority current cross-links are [A1], [A2], [A3], [A4], [A5], [A6], [A7]. These are links among separately custodied works, not silent identity mergers. Their role is to expose mutual dependencies, terminology collisions, and places where a reader should traverse the corpus before treating this paper as self-sufficient.

Reference-pass counterevidence and limit search

The pass applied the following pressure tests: Structural and computational correspondences remain insufficient as phenomenal evidence unless linked to

discriminating empirical predictions against competing theories. Formal vocabularies and graph links make commitments explicit but cannot by themselves establish correct identity, shared interpretation, or lossless translation between contexts. A complete provenance graph can faithfully record a false or unauthorized process; provenance is evidence about lineage, not automatic truth or legitimacy. These limits are not generic disclaimers: they identify the assumptions that a later empirical, formal, or implementation review must try to break.

Reference-pass influence and novelty boundary

No newly added external source is asserted to be a documented authorial influence unless the predecessor source already said so. `Predecessor` means earlier relevant work;

`adjacent` means a useful comparison; `standard` means an interoperability or assurance baseline; and `counterevidence` marks work that narrows or pressures a claim. Novelty therefore remains a reviewable, claim-level question rather than an inference from a low similarity score or an absent citation.

References

External scholarship, standards, and prior art

[**E1**] David J. Chalmers. "Facing Up to the Problem of Consciousness." *Journal of Consciousness Studies* 2(3), 1995, pp. 200-219. https://consc.net/papers/facing.html **Relationship:** Foundational framing of the explanatory gap; cited to bound, not resolve, phenomenal claims.

[**E2**] Larissa Albantakis et al. "Integrated Information Theory (IIT) 4.0: Formulating the Properties of Phenomenal Existence in Physical Terms." *PLOS Computational Biology* 19(10), 2023, e1011465. DOI: 10.1371/journal.pcbi.1011465. https://doi.org/10.1371/journal.pcbi.1011465 **Relationship:** Competing formal theory of consciousness included for explicit contrast and counterevidence obligations.

[**E3**] Stanislas Dehaene and Jean-Pierre Changeux. "Experimental and Theoretical Approaches to Conscious Processing." *Neuron* 70(2), 2011, pp. 200-227. DOI: 10.1016/j.neuron.2011.03.018. https://doi.org/10.1016/j.neuron.2011.03.018 **Relationship:** Competing global-workspace account included to prevent a corpus-native ontology from substituting for empirical neuroscience.

[**E4**] Nikolaus Kriegeskorte, Marieke Mur, and Peter A. Bandettini. "Representational Similarity Analysis - Connecting the Branches of Systems Neuroscience." *Frontiers in Systems Neuroscience* 2, 2008, 4. DOI: 10.3389/neuro.06.004.2008. https://doi.org/10.3389/neuro.06.004.2008 **Relationship:** Predecessor empirical method for comparing representational geometries.

[**E5**] Francesco Locatello et al. "Challenging Common Assumptions in the

Unsupervised Learning of Disentangled Representations." *PMLR* 97, 2019, pp. 4114-4124. https://proceedings.mlr.press/v97/locatello19a.html **Relationship:** Counterevidence showing that desired latent factorization is not identifiable from unsupervised data without inductive biases.

[**E6**] Richard Cyganiak, David Wood, and Markus Lanthaler, eds. *RDF 1.1 Concepts and Abstract Syntax*. W3C Recommendation, 25 February 2014. https://www.w3.org/TR/rdf11-concepts/ **Relationship:** Adjacent graph data model and identity baseline.

[**E7**] Alistair Miles and Sean Bechhofer, eds. *SKOS Simple Knowledge Organization System Reference*. W3C Recommendation, 18 August 2009. https://www.w3.org/TR/skos-reference **Relationship:** Adjacent standard for concepts, labels, schemes, notes, and semantic mappings.

Ableman corpus

[**A1**] Adam Ableman. *Ableman Consciousness Ontology (ACO)*. Ableman Corpus current edition; work identity `PAPER.PUB.18203749`;
source SHA-256
`ACD97032780DA7B3A36750702BE1F47BCB2E6361EF6E74ECA8B441D7C6A25DDE`.
Corpus relationship: priority cross-link selected by corpus profile and similarity review; no identity merger or publication effect.

[**A2**] Adam Ableman. *The Locus of Qualia*. Ableman Corpus current edition; work identity
`AMCS-LOCAL-WORK-AF27E66B81811DB1`;
source SHA-256
`7B83EBD0FFDCFEBFA522A715AE0E4D1049FCA0B4CE0618C2017F8BB9206E5424`.
Corpus relationship: priority cross-link selected by corpus profile and similarity review; no identity merger or publication effect.

[**A3**] Adam Ableman. *Possible Minds: Against the Provincialization of Consciousness*. Ableman Corpus current edition; work identity `PAPER.PUB.20681869`;
source SHA-256
`77F381A29E2EF801085938DB369F3F432321CCEB8B9DD3969330337C24BA4998`.
Corpus relationship: priority cross-link selected by corpus profile and similarity review; no identity merger or publication effect.

[**A4**] Adam Ableman. *Admissibility Theory of Consciousness*. Ableman Corpus current edition; work identity
`AMCS-LOCAL-WORK-8FDC994841DA6CF3`;
source SHA-256
`F409652555CEC4898DBEC1E8BBB179CEE020E7FF35F4C8F307CDA82253D6F49F`.
Corpus relationship: priority cross-link selected by corpus profile and similarity review; no identity merger or publication effect.

[**A5**] Adam Ableman. *Meaning as Binding*. Ableman Corpus current edition;

work identity
AMCS-LOCAL-WORK-5F976917F0B283EA;
source SHA-256
A470D1A0B9ED7C3CEAB6BB88BE4B642A2BA36427447F787764329FD574DA0C97.
Corpus relationship: priority cross-link selected by corpus profile and similarity review; no identity merger or publication effect.

[**A6**] Adam Ableman. *Governed Meaning Base (GMB).* Ableman Corpus current edition; work identity
AMCS-LOCAL-WORK-D95CCC4FB7B58BFB;
source SHA-256
A3B62E50F8BEE125FDAE0D38C2839AFD1CCC138296268BF93217481E38AA48EA.
Corpus relationship: priority cross-link selected by corpus profile and similarity review; no identity merger or publication effect.

[**A7**] Adam Ableman. *Semantic Manifold Topology.* Ableman Corpus current edition; work identity PAPER.PUB.20257244;
source SHA-256
B8F28CC652A73A3EE4CF6FD0E23C1ABC922A6660319B2FB4CC1CCAAA606EC64A.
Corpus relationship: priority cross-link selected by corpus profile and similarity review; no identity merger or publication effect.

Editorial source notes retained from predecessor

Ableman Corpus References

Adam Ableman. *A Formal Theory of Work Custody.* Ableman Corpus current edition; conceptual claims retain the standing stated in the work.

Adam Ableman. *The Constitutional Observability Platform.* Ableman Corpus current edition; conceptual claims retain the standing stated in the work.

Adam Ableman. *Governed Context & Memory (GCM): Context Configuration as Constitutional Surface, Semantic Hashing as Enforcement, and Explicit Ignorance as First-Class Success in Generational AI.* Zenodo, 2026. https://doi.org/10.5281/zenodo.18633797.

Adam Ableman. *Governed Meaning Base (GMB).* Ableman Corpus current edition in the author-confirmed historical terminology CCF -> CMB -> GMB; conceptual claims retain the standing stated in the work.

Adam Ableman. *Meaning as Binding.* Ableman Corpus current edition; conceptual claims retain the standing stated in the work.

Adam Ableman. *Reasoning by Way of Receipts.* Ableman Corpus current edition; conceptual claims retain the standing stated in the work.

Adam Ableman. *Semantic Custody.* Zenodo, 2026. https://doi.org/10.5281/zenodo.20300952.

Adam Ableman. *Semantic Legitimacy Before Context.* Ableman Corpus current edition; conceptual claims retain the standing stated in the work.

Adam Ableman. *The Three Hash Rule*. Zenodo, 2026. https://doi.org/10.5281/zenodo.20788617.

Biocca, F. (1997). The cyborg's dilemma: Progressive embodiment in virtual environments. *Journal of Computer-Mediated Communication, 3*(2).

Clark, A. (1997). *Being there: Putting brain, body, and world together again*. MIT Press.

Clark, A., & Chalmers, D. J. (1998). The extended mind. *Analysis, 58*(1), 7–19.

Dennett, D. C. (1991). *Consciousness explained*. Little, Brown and Company.

Gibson, J. J. (1979). *The ecological approach to visual perception*. Houghton Mifflin.

Heeter, C. (1992). Being there: The subjective experience of presence. *Presence: Teleoperators and Virtual Environments, 1*(2), 262–271.

Lombard, M., & Ditton, T. (1997). At the heart of it all: The concept of presence. *Journal of Computer-Mediated Communication, 3*(2).

Merleau-Ponty, M. (2012). *Phenomenology of perception* (D. A. Landes, Trans.). Routledge. (Original work published 1945)

Metzinger, T. (2003). *Being no one: The self-model theory of subjectivity*. MIT Press.

Murray, J. H. (1997). *Hamlet on the holodeck: The future of narrative in cyberspace*. Free Press.

Noe, A. (2004). *Action in perception*. MIT Press.

Reeves, B., & Nass, C. (1996). *The media equation: How people treat computers, television, and new media like real people and places*. Cambridge University Press.

Slater, M. (2009). Place illusion and plausibility can lead to realistic behaviour in immersive virtual environments. *Philosophical Transactions of the Royal Society B: Biological Sciences, 364*(1535), 3549–3557.

Weizenbaum, J. (1966). ELIZA–A computer program for the study of natural-language communication between man and machine. *Communications of the ACM, 9*(1), 36–45.

The external reference list is retained from the author source. Current source verification and literature review remain required across consciousness studies, phenomenology, presence, anthropomorphism, human-computer interaction, persuasive and adaptive systems, human-AI attachment, epistemic calibration, psychiatry, mania, psychosis, dissociation, altered states, sleep deprivation, and research ethics for a future evidence-qualified revision. The current list does not by itself support the paper's newer AI- and clinical-risk propositions.

The entries below are corpus-internal comparison and dependency sources. Their inclusion does not establish external validation of this paper's empirical, safety, novelty, or performance claims.

Adam Ableman *The Three Hash Rule.* Zenodo, 2026. doi:10.5281/zenodo.20788617

Adam Ableman *Governed Context & Memory (GCM): Context Configuration as Constitutional Surface, Semantic Hashing as Enforcement, and Explicit Ignorance as First-Class Success in Generational AI.* Zenodo, 2026. doi:10.5281/zenodo.18633797

Adam Ableman *Semantic Custody.* Zenodo, 2026. doi:10.5281/zenodo.20300952

Adam Ableman *A Theory of Agentic Observation.* Zenodo, 2026. doi:10.5281/zenodo.20437813

Adam Ableman *CARA: Constrained Artificial Reasoning Architecture.* Zenodo, 2025. doi:10.5281/zenodo.18099715

Adam Ableman *Constitutional Computation.* Zenodo, 2026. doi:10.5281/zenodo.20624488

Adam Ableman *Constraint-Bearing Composability.* Zenodo, 2026. doi:10.5281/zenodo.21632061

Adam Ableman *Fractal Governance Architecture.* Zenodo, 2026. doi:10.5281/zenodo.20754189

Corpus source and companion formats: https://github.com/AblemanMosaic/ableman-corpus
License: CC-BY-ND-4.0

DOI: 10.5281/zenodo.23848465